

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import warnings
6 warnings.filterwarnings('ignore')
```

In [2]:

```
1 df=pd.read_csv("cars.csv")
```

In [3]:

```
1 df.head()
```

Out[3]:

	symboling	normalized- losses	make	fuel- type	body- style	drive- wheels	engine- location	width	height	engine- type
0	3	?	alfa- romero	gas	convertible	rwd	front	64.1	48.8	dohc
1	3	?	alfa- romero	gas	convertible	rwd	front	64.1	48.8	dohc
2	1	?	alfa- romero	gas	hatchback	rwd	front	65.5	52.4	ohcv
3	2	164	audi	gas	sedan	fwd	front	66.2	54.3	ohc
4	2	164	audi	gas	sedan	4wd	front	66.4	54.3	ohc



In [4]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              205 non-null    int64
1   normalized-losses      205 non-null    object
2   make                   205 non-null    object
3   fuel-type              205 non-null    object
4   body-style             205 non-null    object
5   drive-wheels           205 non-null    object
6   engine-location        205 non-null    object
7   width                  205 non-null    float64
8   height                 205 non-null    float64
9   engine-type            205 non-null    object
10  engine-size            205 non-null    int64
11  horsepower              205 non-null    object
12  city-mpg                205 non-null    int64
13  highway-mpg             205 non-null    int64
14  price                   205 non-null    int64
dtypes: float64(2), int64(5), object(8)
memory usage: 24.1+ KB
```

In [5]:

```
1 # total no of records :205, column 15
2 # normalized-losses and horsepower contains numeric data but they are in object .
```

In [6]:

```
1 df['normalized-losses'].value_counts()
```

Out[6]:

?	41
161	11
91	8
150	7
134	6
128	6
104	6
85	5
94	5
65	5
102	5
74	5
168	5
103	5
95	5
106	4
93	4
118	4
148	4
122	4
83	3
125	3
154	3
115	3
137	3
101	3
119	2
87	2
89	2
192	2
197	2
158	2
81	2
188	2
194	2
153	2
129	2
108	2
110	2
164	2
145	2
113	2
256	1
107	1
90	1
231	1
142	1
121	1
78	1
98	1
186	1
77	1

Name: normalized-losses, dtype: int64

In [7]:

```
1 df['horsepower'].value_counts()
```

Out[7]:

68	19
70	11
69	10
116	9
110	8
95	7
88	6
62	6
101	6
160	6
114	6
84	5
97	5
102	5
145	5
82	5
76	5
111	4
92	4
123	4
86	4
90	3
73	3
85	3
207	3
182	3
121	3
152	3
112	2
56	2
161	2
156	2
94	2
52	2
?	2
162	2
155	2
184	2
100	2
176	2
55	1
262	1
134	1
115	1
140	1
48	1
58	1
60	1
78	1
135	1
200	1
64	1
120	1
72	1

```
154      1
288      1
143      1
142      1
175      1
106      1
Name: horsepower, dtype: int64
```

In [8]:

```
1 df['normalized-losses'].replace("?", np.nan, inplace=True)
```

In [9]:

```
1 df.isnull().sum()
```

Out[9]:

```
symboling      0
normalized-losses  41
make           0
fuel-type      0
body-style     0
drive-wheels   0
engine-location 0
width          0
height         0
engine-type    0
engine-size    0
horsepower     0
city-mpg       0
highway-mpg    0
price          0
dtype: int64
```

In [10]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              205 non-null    int64
1   normalized-losses      164 non-null    object
2   make                   205 non-null    object
3   fuel-type              205 non-null    object
4   body-style             205 non-null    object
5   drive-wheels           205 non-null    object
6   engine-location        205 non-null    object
7   width                  205 non-null    float64
8   height                 205 non-null    float64
9   engine-type            205 non-null    object
10  engine-size            205 non-null    int64
11  horsepower             205 non-null    object
12  city-mpg               205 non-null    int64
13  highway-mpg            205 non-null    int64
14  price                  205 non-null    int64
dtypes: float64(2), int64(5), object(8)
memory usage: 24.1+ KB
```

In [11]:

```
1 df['normalized-losses']=df['normalized-losses'].astype('float64')
```

In [12]:

```
1 n_mean=df['normalized-losses'].mean()
```

In [13]:

```
1 n_mean
```

Out[13]:

122.0

In [14]:

```
1 df['normalized-losses'].fillna(n_mean,inplace=True)
```

In [15]:

```
1 df.isnull().sum()
```

Out[15]:

```
symboling          0
normalized-losses  0
make              0
fuel-type         0
body-style        0
drive-wheels      0
engine-location   0
width            0
height           0
engine-type       0
engine-size       0
horsepower        0
city-mpg          0
highway-mpg       0
price            0
dtype: int64
```

In [16]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   symboling             205 non-null   int64
 1   normalized-losses     205 non-null   float64
 2   make                  205 non-null   object
 3   fuel-type             205 non-null   object
 4   body-style            205 non-null   object
 5   drive-wheels          205 non-null   object
 6   engine-location       205 non-null   object
 7   width                 205 non-null   float64
 8   height                205 non-null   float64
 9   engine-type           205 non-null   object
10   engine-size           205 non-null   int64
11   horsepower            205 non-null   object
12   city-mpg              205 non-null   int64
13   highway-mpg           205 non-null   int64
14   price                 205 non-null   int64
dtypes: float64(3), int64(5), object(7)
memory usage: 24.1+ KB
```

In [17]:

```
1 df.head(10)
```

Out[17]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type
0	3	122.0	alfa-romero	gas	convertible	rwd	front	64.1	48.8	dohc
1	3	122.0	alfa-romero	gas	convertible	rwd	front	64.1	48.8	dohc
2	1	122.0	alfa-romero	gas	hatchback	rwd	front	65.5	52.4	ohcv
3	2	164.0	audi	gas	sedan	fwd	front	66.2	54.3	ohc
4	2	164.0	audi	gas	sedan	4wd	front	66.4	54.3	ohc
5	2	122.0	audi	gas	sedan	fwd	front	66.3	53.1	ohc
6	1	158.0	audi	gas	sedan	fwd	front	71.4	55.7	ohc
7	1	122.0	audi	gas	wagon	fwd	front	71.4	55.7	ohc
8	1	158.0	audi	gas	sedan	fwd	front	71.4	55.9	ohc
9	0	122.0	audi	gas	hatchback	4wd	front	67.9	52.0	ohc

In [18]:

```
1 df['horsepower'].replace("?", np.nan, inplace=True)
```

In [19]:

```
1 df.isnull().sum()
```

Out[19]:

```
symboling          0
normalized-losses  0
make              0
fuel-type          0
body-style         0
drive-wheels       0
engine-location    0
width             0
height            0
engine-type        0
engine-size        0
horsepower         2
city-mpg           0
highway-mpg        0
price             0
dtype: int64
```



In [20]:

```
1 df.dropna(inplace=True)
```

In [21]:

```
1 df.isnull().sum()
```

Out[21]:

```
symboling          0
normalized-losses  0
make              0
fuel-type         0
body-style        0
drive-wheels      0
engine-location   0
width            0
height           0
engine-type       0
engine-size       0
horsepower        0
city-mpg          0
highway-mpg       0
price            0
dtype: int64
```

In [22]:

```
1 df['horsepower'] = df['horsepower'].astype('float64')
```

In [23]:

```
1 df.dtypes
```

Out[23]:

```
symboling          int64
normalized-losses  float64
make              object
fuel-type         object
body-style        object
drive-wheels      object
engine-location   object
width            float64
height           float64
engine-type       object
engine-size       int64
horsepower        float64
city-mpg          int64
highway-mpg       int64
price            int64
dtype: object
```

In [24]:


```
1 #handling outlier
```

In [25]:

```
1 plt.figure(figsize=(8,8))
2 sns.boxplot(data=df,x="price",y="make")
```

Out[25]:

<AxesSubplot:xlabel='price', ylabel='make'>



In [26]:

```
1 df.loc[(df['fuel-type']=="gas")]
```

Out[26]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type
0	3	122.0	alfa-romero	gas	convertible	rwd	front	64.1	48.8	dohc
1	3	122.0	alfa-romero	gas	convertible	rwd	front	64.1	48.8	dohc
2	1	122.0	alfa-romero	gas	hatchback	rwd	front	65.5	52.4	ohcv
3	2	164.0	audi	gas	sedan	fwd	front	66.2	54.3	ohc
4	2	164.0	audi	gas	sedan	4wd	front	66.4	54.3	ohc
...	...	...	...	...	...	...	...	...	...	...
199	-1	74.0	volvo	gas	wagon	rwd	front	67.2	57.5	ohc
200	-1	95.0	volvo	gas	sedan	rwd	front	68.9	55.5	ohc
201	-1	95.0	volvo	gas	sedan	rwd	front	68.8	55.5	ohc
202	-1	95.0	volvo	gas	sedan	rwd	front	68.9	55.5	ohcv
204	-1	95.0	volvo	gas	sedan	rwd	front	68.9	55.5	ohc

183 rows × 11 columns

In [27]:

```
1 df.loc[(df['make']=='toyota') & (df['price']>14000)]
```

Out[27]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type
172	2	134.0	toyota	gas	convertible	rwd	front	65.6	53.0	ohc
178	3	197.0	toyota	gas	hatchback	rwd	front	67.7	52.0	dohc
179	3	197.0	toyota	gas	hatchback	rwd	front	67.7	52.0	dohc
180	-1	90.0	toyota	gas	sedan	rwd	front	66.5	54.1	dohc
181	-1	122.0	toyota	gas	wagon	rwd	front	66.5	54.1	dohc

In [28]:


```
1 df.drop([172,178,179,180,181],axis=0,inplace=True)
```

In [40]:

```
1 plt.figure(figsize=(9,5))
2 sns.boxplot(data=df,x="price",y="make")
```

Out[40]:

<AxesSubplot:xlabel='price', ylabel='make'>



In [30]:

```
1 df.loc[(df['make']=='mitsubishi') & (df['price']>11000)]
```

Out[30]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine type
82	3	122.0	mitsubishi	gas	hatchback	fwd	front	66.3	50.2	oh
83	3	122.0	mitsubishi	gas	hatchback	fwd	front	66.3	50.2	oh
84	3	122.0	mitsubishi	gas	hatchback	fwd	front	66.3	50.2	oh

In [31]:

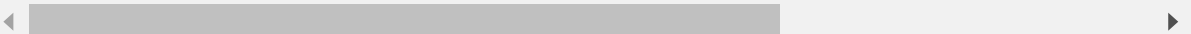
```
1 df.drop([82,83,84],axis=0,inplace=True)
```

In [32]:

```
1 df.loc[(df['make']=='plymouth')& (df['price']>10000)]
```

Out[32]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type
124	3	122.0	plymouth	gas	hatchback	rwd	front	66.3	50.2	ohc



In [33]:

```
1 df.loc[(df['make']=='isuzu')& (df['price']>20000)]
```

Out[33]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type	engine-size
45	0	122.0	isuzu	gas	sedan	fwd	front	63.6	52.0	ohc	



In [34]:

```
1 df.loc[(df['make']=='honda')& (df['price']>11000)]
```

Out[34]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type	engine-size
41	0	85.0	honda	gas	sedan	fwd	front	65.2	54.1	ohc	

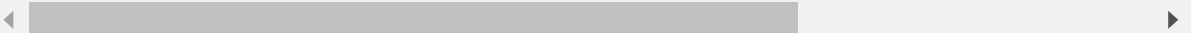


In [35]:

```
1 df.loc[(df['make']=='dodge')& (df['price']>10000)]
```

Out[35]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type	engine-size
29	3	145.0	dodge	gas	hatchback	fwd	front	66.3	50.2	ohc	



In [36]:


```
1 df.drop([124,45,41,29],axis=0,inplace=True)
```

In [42]:

```
1 plt.figure(figsize=(9,5))
2 sns.boxplot(data=df,x="price",y="make")
```

Out[42]:

<AxesSubplot:xlabel='price', ylabel='make'>



In [ ]:

1